Mir Sufiyan Ali

Hyderabad, India

+91 9390822459 | mirsufiyanali6s@gmail.com | LinkedIn | GitHub | Portfolio

Professional Summary

Data Analyst with hands-on experience in SQL Server (T-SQL), Python, and Power BI, turning raw transactional data into validated, decision-ready insights. Grounds every finding in statistical testing before it reaches a dashboard, and has delivered end-to-end analytics across finance, healthcare, E-commerce, and retail domains that directly informed business decisions.

Education

Lords Institute of Engineering and Technology — 2023 – 2027

Bachelor of Engineering in Computer Science (Data Science) | CGPA: 9.02/10 — *Hyderabad, India*

Relevant Coursework

- Business Analytics
- Machine Learning
- Data Modeling
- DBMS
- Exploratory Data Analysis
- Probability & Statistics
- Data Visualization
- Data Warehousing

Experience

Manac Infotech — Jul – Aug 2025 | GitHub

Data Science Intern — *India*

- Built a Python/Scikit-Learn ML pipeline (SVM, KNN) to detect water-meter tampering and non-technical losses from billing data, cutting false-positive tamper alerts from 35% to 12%.
- Engineered 22 behavioral features (usage trend slope, rolling deltas, deviation ratios, tamper flags) from raw consumption data to generate calibrated fraud-risk scores per account.
- Developed a Django + SQL Server dashboard (ChartJS) to rank high-risk accounts for field audits and validated flagged accounts against confirmed theft cases, presenting results to stakeholders.

Projects

VendorIQ & Inventory Analytics | *SQL Server (T-SQL), Python, Power BI* — GitHub

- Modeled 6 relational tables (15.6M rows, 129 vendors, 268K products) into a star schema and ran 48 SQL business queries, finding 15 vendors (12% of the base) drive 37.73% of revenue while 7 contribute just 0.02% — informing a vendor-base consolidation recommendation.
- Statistically validated that bulk orders (10+ units) cut unit cost 32.6% ($9.85 vs. $14.62, Mann–Whitney U, $p < 0.001$) and flagged $44.85M tied up in 901 high-value outlier SKUs (9.33% of catalog), directing bulk-purchasing and inventory-reduction decisions in a 3-page Power BI dashboard.

Healthcare Operations & Financial Performance Analytics | *SQL Server, Python, Power BI* — GitHub

- Cleaned and standardized 54,966 hospital admission records (534 duplicates removed) across 28 SQL business queries; used t-test, ANOVA, and chi-square testing to rule out age, admission type, and gender as significant drivers of billing or length of stay, preventing leadership from chasing false cost signals.
- Caught 106 records with negative billing (down to -$2,008 each) as a data-quality defect and flagged near-uniform hospital/doctor distributions as evidence the dataset was synthetic — recommending the $1.40B billing dashboard be read as directional, not a real hospital leaderboard.

Banking Intelligence: Customer & Transaction Analytics | *SQL Server, Python, Power BI* — GitHub

- Surfaced that 67% of the customer base is Suspended or Inactive against only 33% Active — the central retention risk in a $25.08M transaction book — then used t-test, ANOVA, and chi-square testing to rule out debit/credit split, channel, and customer status as the cause.
- Found transaction frequency near-perfectly predicts customer value ($r = 0.99$) while age does not ($r \approx 0.06$), and identified ATM as the highest-failure channel (5.34% vs. 3.95% web) across 18 SQL queries — informing a frequency-based retention target and a channel-reliability fix.

Seller Concentration & Freight Risk Analysis | *SQL Server, Python, Power BI* — GitHub

- Modeled 9 relational tables across a 100K-order marketplace and found the top 10% of sellers generate 67.5% of revenue (top 10 alone: 13.15%) while revenue growth had turned from a steady 2017 climb into a -4.56% contraction — flagging seller concentration as a near-term risk, not just a structural one.
- Confirmed cross-state shipments cost significantly more in freight than same-state (Welch's t-test, $p < 0.001$), consuming 16.57% of item revenue, and found only 1.0% of orders span more than one category despite repeat cross-category buyers — directing freight-pricing and cross-sell recommendations across 5 dashboards.

Technical Skills

Programming & Querying: SQL, T-SQL, Python

Data Analysis & ML: EDA, Statistical Analysis, Feature Engineering, Data Cleaning, Anomaly Detection, Customer Segmentation, Pandas

Visualization & BI: Power BI, Data Modeling, Dashboard Development, KPI Reporting, Excel (Pivot Tables, VLOOKUP), Matplotlib, Business Intelligence

Databases, Tools & Concepts: SQL Server (SSMS), Data Modeling (Star Schema), Git, GitHub, Jupyter Notebook, Google Colab, Stakeholder Communication, Requirements Gathering, Data Storytelling

Certifications

Google Advanced Data Analytics Professional Certificate — Apr – Aug 2026 | Certificate

Google